

## Phase\_\_2\_\_CF

November 30, 2019

### 0.0.1 Type III: Collaborative Filtering based Recommendation System

It works on the basic idea that people with similar characteristics have similar interests in movies.

There are 2 types of Collaborative Filtering Techniques:

*User-Based:* Here we assume that similar people have similar liking. If user A and user B have watched a lot of similar movies and have rated them almost equally, then it can be inferred that their interests match. Thus we can recommend movies that A have liked to B.

*Item-Based:* Here we take into consideration the content of movies as well and not only user similarity. Here we find those movies which the user would enjoy based on his previously rated movies.

User Based Collaborative Filtering is a very common and highly used approach for movielens dataset because of the fact that:

- It is easy to implement.
- It is context independent.
- It gives better results than content based recommendation system.

But this approach is not the best. Recommending movies without taking into consideration the content of the movie would never give best results.

Thus we tried to implement an **item based recommendation system** with our movies dataset.

Firstly, we start off by reading the data. To find similarity between items, we are using 5 different types of feature combinations:

1. Genre, Crew
2. Cast, Keywords
3. Genre, Crew, Production Countries
4. Genre, Production Countries, Keywords
5. Genre, Crew, Production Countries, Keywords

The above combinations were made on our basic understanding of what all features influence our liking for a movie.

The data is combined from multiple files: keywords.csv, credits.csv, movies\_metadata.csv, ratings.csv

### Mount & Imports

```
[0]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
[0]: cd /content/drive/My\ Drive/DataScience\ Project/the-movies-dataset
```

/content/drive/My Drive/DataScience Project/the-movies-dataset

```
[0]: ls
```

```
BestFeatureVector.csv    Errors1.csv    links_small.csv
cosSimScore.csv          Errors-S.csv   movies_metadata.csv
credits.csv              keywords.csv   ratings.csv
'Customer Feedback.csv'  links.csv      ratings_small.csv
```

```
[0]: from sklearn.metrics.pairwise import cosine_similarity
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import MultiLabelBinarizer
from nltk.corpus import stopwords
from ast import literal_eval
import pandas as pd, numpy as np, matplotlib.pyplot as plt
import nltk, re, string
from prettytable import PrettyTable
from prettytable import from_csv
nltk.download('stopwords')
stopWord = set(stopwords.words('english'))
```

[nltk\_data] Downloading package stopwords to /root/nltk\_data...

[nltk\_data] Package stopwords is already up-to-date!

## Functions

```
[0]: def removeStopWord(tokens): #Function removing stop words and punctuation
    global stopWord
    low = tokens.lower()
    cl = re.compile('_') #remove punctuation
    loweredTokens = re.split(';|_|!|#|@|$|%|\^|&|\*|\(|\)|:|\\"|\'| |\\.|\\.|,|', low)
    fr = []
    for i in loweredTokens:
        if(i not in stopWord):
            fr.append(i)
    return (' ').join(fr)
```

```
[0]: def binarise(y_train): #Function to create a binary vector for features
    mlb = MultiLabelBinarizer()
    y_Bin = mlb.fit_transform(y_train)
    clsNames = mlb.classes_
    return y_Bin, clsNames
```

```
[0]: def vectorForm(df,col): #Creates feature vector
    temp = [[] for i in range(len(df))]
    colNames = []
    for i in col:
        val,colName = binarise(df[i])
        temp = np.hstack((temp,val))
        colNames.extend(colName)
    vectDf = pd.DataFrame(temp,columns=colNames)
    finlVec = pd.concat([vectDf,df.rating],axis=1)
    return finlVec
```

```
[0]: def recommend(test,r,sim,mu,title,n): #Recommends top n movies
    predRating = dict()
    for i in test.index:
        sum1 = 0
        sum2 = 0
        for j in test.index:
            sum1 += (sim[i][j]*(r[j]-mu))
            sum2 += sim[i][j]
        predRating[i] = round(mu + (sum1/sum2),2)
    sortedScore = sorted(predRating.items(), key=lambda x: x[1],reverse=True)
    ↪#Sorting in decreasing order on the basis of similarity score
    topScores = sortedScore[1:n+1]
    movies = [i[0] for i in topScores]
    return list(title[movies])
```

## Reading Data and preprocessing

```
[0]: keyDf = pd.read_csv('keywords.csv')
    credDf = pd.read_csv('credits.csv')
    metaDf = pd.read_csv('movies_metadata.csv')
    ratDf = pd.read_csv('ratings.csv')
```

```
/usr/local/lib/python3.6/dist-packages/IPython/core/interactiveshell.py:2718:
DtypeWarning: Columns (10) have mixed types. Specify dtype option on import or
set low_memory=False.
    interactivity=interactivity, compiler=compiler, result=result)
```

```
[0]: print(keyDf.shape,credDf.shape,metaDf.shape,ratDf.shape)
```

```
(46419, 2) (45476, 3) (45466, 24) (26024289, 4)
```

```
[0]: ratDf.head()
```

```
[0]:   userId  movieId  rating  timestamp
0      1      110     1.0  1425941529
1      1      147     4.5  1425942435
2      1      858     5.0  1425941523
3      1     1221     5.0  1425941546
4      1     1246     5.0  1425941556
```

```
[0]: ratDf.rename(columns = {"movieId": "id"},inplace=True) #Renaming movieId col to id
↳ id for merging
```

```
[0]: ratDf.columns
```

```
[0]: Index(['userId', 'id', 'rating', 'timestamp'], dtype='object')
```

```
[0]: metaDf.head(2)
```

```
[0]:   adult  ... vote_count
0  False  ...    5415.0
1  False  ...    2413.0

[2 rows x 24 columns]
```

```
[0]: keyDf.head()
```

```
[0]:      id      keywords
0    862  [{'id': 931, 'name': 'jealousy'}, {'id': 4290,...
1   8844  [{'id': 10090, 'name': 'board game'}, {'id': 1...
2  15602  [{'id': 1495, 'name': 'fishing'}, {'id': 12392...
3  31357  [{'id': 818, 'name': 'based on novel'}, {'id':...
4  11862  [{'id': 1009, 'name': 'baby'}, {'id': 1599, 'n...
```

```
[0]: credDf.head()
```

```
[0]:      cast  ...  id
0  [{'cast_id': 14, 'character': 'Woody (voice)',...  ...    862
1  [{'cast_id': 1, 'character': 'Alan Parrish', '...  ...    8844
2  [{'cast_id': 2, 'character': 'Max Goldman', 'c...  ...    15602
3  [{'cast_id': 1, 'character': 'Savannah 'Vannah...  ...    31357
4  [{'cast_id': 1, 'character': 'George Banks', '...  ...    11862

[5 rows x 3 columns]
```

Info about DF

```
[0]: print(keyDf.info())
      print()
      print(credDf.info())
      print()
      print(metaDf.info())
      print()
      print(ratDf.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 46419 entries, 0 to 46418
Data columns (total 2 columns):
id          46419 non-null int64
keywords    46419 non-null object
dtypes: int64(1), object(1)
memory usage: 725.4+ KB
None
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45476 entries, 0 to 45475
Data columns (total 3 columns):
cast        45476 non-null object
crew         45476 non-null object
id           45476 non-null int64
dtypes: int64(1), object(2)
memory usage: 1.0+ MB
None
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45466 entries, 0 to 45465
Data columns (total 24 columns):
adult                45466 non-null object
belongs_to_collection  4494 non-null object
budget               45466 non-null object
genres               45466 non-null object
homepage             7782 non-null object
id                   45466 non-null object
imdb_id              45449 non-null object
original_language     45455 non-null object
original_title        45466 non-null object
overview             44512 non-null object
popularity            45461 non-null object
poster_path          45080 non-null object
production_companies  45463 non-null object
production_countries  45463 non-null object
release_date          45379 non-null object
revenue               45460 non-null float64
runtime              45203 non-null float64
```

```

spoken_languages      45460 non-null object
status                45379 non-null object
tagline               20412 non-null object
title                 45460 non-null object
video                 45460 non-null object
vote_average          45460 non-null float64
vote_count            45460 non-null float64
dtypes: float64(4), object(20)
memory usage: 8.3+ MB
None

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26024289 entries, 0 to 26024288
Data columns (total 4 columns):
userId      int64
id           int64
rating       float64
timestamp    int64
dtypes: float64(1), int64(3)
memory usage: 794.2 MB
None

```

Converting int to str

```

[0]: keyDf.id = keyDf.id.apply(str)
      credDf.id = credDf.id.apply(str)
      ratDf.id = ratDf.id.apply(str)

```

Sampling of data

```

[0]: maxUserRat = list(ratDf.userId[ratDf.userId == ratDf.userId.value_counts().
      ↪index[0]].index) #Picking users with maximum ratings count i.e. 45811
      sampleRatDf = ratDf.iloc[maxUserRat,:] #Sampling maximum users data based on
      ↪ratings
      sampleRatDf.drop(['userId','timestamp'],axis=1,inplace=True) #Dropping off
      ↪timestamp & userId
      sampleRatDf.index = range(len(sampleRatDf)) #Resetting indexes
      sampleRatDf.head()

```

```

/usr/local/lib/python3.6/dist-packages/pandas/core/frame.py:4117:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

```

See the caveats in the documentation: [http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
errors=errors,

```
[0]:   id  rating
      0    1    4.0
      1    2    3.0
      2    6    4.0
      3    9    2.5
      4   10    3.0
```

Merging on movieid column

```
[0]: merKeyDf = metaDf.merge(keyDf,on='id')
      merCredDf = merKeyDf.merge(credDf,on='id')
      mergedDf = merCredDf.merge(sampleRatDf,on='id')
```

```
[0]: mergedDf.shape
```

```
[0]: (3634, 28)
```

```
[0]: finalDf =
      ↪mergedDf[['id','title','genres','keywords','cast','crew','production_countries','rating']]
      finalDf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3634 entries, 0 to 3633
Data columns (total 8 columns):
id                3634 non-null object
title             3634 non-null object
genres            3634 non-null object
keywords          3634 non-null object
cast              3634 non-null object
crew              3634 non-null object
production_countries 3634 non-null object
rating            3634 non-null float64
dtypes: float64(1), object(7)
memory usage: 255.5+ KB
```

Preprocessing (Converting structures to list & removing stop words or punctuation)

```
[0]: finalDf['genres'] = finalDf['genres'].fillna('').apply(literal_eval).
      ↪apply(lambda x: [removeStopWord(i['name']) for i in x] if isinstance(x,
      ↪list) else [])
      finalDf['cast'] = finalDf['cast'].fillna('').apply(literal_eval).apply(lambda
      ↪x: [removeStopWord(i['name']) for i in x] if isinstance(x, list) else [])
      finalDf['crew'] = finalDf['crew'].fillna('').apply(literal_eval).apply(lambda
      ↪x: [removeStopWord(i['name']) for i in x] if isinstance(x, list) else [])
      finalDf['production_countries'] = finalDf['production_countries'].fillna('').
      ↪apply(literal_eval).apply(lambda x: [removeStopWord(i['name']) for i in x]
      ↪if isinstance(x, list) else [])
```

```
finalDf['keywords'] = finalDf['keywords'].fillna('').apply(literal_eval).
↳ apply(lambda x: [removeStopWord(i['name']) for i in x] if isinstance(x, list)
↳ else [])
```

/usr/lib/python3.6/re.py:212: FutureWarning: split() requires a non-empty pattern match.

```
return _compile(pattern, flags).split(string, maxsplit)
```

/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:1:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
"""Entry point for launching an IPython kernel.
```

/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:2:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:3:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

This is separate from the ipykernel package so we can avoid doing imports until

/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:4:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
after removing the cwd from sys.path.
```

/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:5:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [http://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)



```
"""
```

```
[0]: finalDf.head()
```

```
[0]:      id  ... rating
0  8844  ...    4.0
1   949  ...    4.0
2  1408  ...    3.5
3   524  ...    4.0
4  8012  ...    3.5
```

```
[5 rows x 8 columns]
```

**Building Recommendation System** The recommendation system is build by following following steps:

Following set of feature combinations is used: 1. Genre, Crew 2. Crew, Keywords 3. Genre ,Crew ,Production Countries 4. Genre, Production Countries, Keywords 5. Genre, Crew, Production Countries, Keywords

Below Steps are followed for all combinations of features: 1. Given dataset for a specific user is split into train - test ratio. 2. Using the training data, where we have user rating corresponding to each movie vector, we train the regressor to predict user's ratings on movies from the test dataset. 3. Now we have predicted user rating for movie i and average rating of movie by all the users. 4. Using RMSE and MAE, we find the est set of features which is able to generalise the user's interests.

Using the best set of features, we find cosine similarity between movies features. Using the below mentioned formula, we find rating of movie for a user.

Using predicted rating for each movie, we sort the best rated movies for user 'u' and recommend.

```
[0]: '''Dictionary for few combination of features'''
```

```
feat = {0:['genres','crew'],
        1:['Crew','keywords'],
        2:['genres','crew','production_countries'],
        3:['genres','production_countries','keywords'],
        4:['genres','crew','production_countries','keywords']}
}
```

```
[0]: '''Running for all combination and calculating MAE and RMSE'''
```

```
testrmse = []
testmae = []
trainrmse = []
trainmae = []
j=0
for i in feat.values():
    j+=1
```

```

# print('\nFor: ',i)
vector = vectorForm(finalDf,i)
vector = vector.values
X = vector[:, :-1]
y = vector[:, -1]
y = y.astype(float)
print(X.shape)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
↪2,shuffle=False)
X_tr, X_ts, y_tr, y_ts = train_test_split(X_train, y_train, test_size=0.
↪2,shuffle=False)
clftst = KNeighborsRegressor(n_neighbors=5)
clftrn = KNeighborsRegressor(n_neighbors=5)
clftst.fit(X_train, y_train)
clftrn.fit(X_tr, y_tr)
predTest = clftst.predict(X_test)
predTrain = clftrn.predict(X_ts)
trm = mean_absolute_error(y_ts, predTrain)
trr = mean_squared_error(y_ts, predTrain)
tsm = mean_absolute_error(y_test, predTest)
tsr = mean_squared_error(y_test, predTest)
testrmse.append(np.sqrt(tsr))
testmae.append(tsm)
trainrmse.append(np.sqrt(trr))
trainmae.append(trm)
x.add_row([str(j), i, X.shape, trm, np.sqrt(trr), tsm, np.sqrt(tsr) ])

```

```

[0]: scoreDf.to_csv('Errors.csv', index=None) #Dumping errors in csv

```

```

[0]: x = [1,2,3,4,5]
y1b1 = [0.0,0.2,0.4,0.6,0.8,1.0]
error = pd.read_csv('Errors1.csv')

fig, axs = plt.subplots(2, 2)
axs[0, 0].bar(x,error['Train MAE'])
axs[0, 0].set_title('MAE of 80% Train Data')
axs[0, 1].bar(x,error['Test MAE'])
axs[0, 1].set_title('MAE of 20% Test Data')
axs[1, 0].bar(x,error['Train RMSE'])
axs[1, 0].set_title('RMSE of 80% Train Data')
axs[1, 1].bar(x,error['Test RMSE'])
axs[1, 1].set_title('RMSE of 20% Test Data')

for ax in axs.flat:
    ax.set(xlabel='Feature Set', ylabel='Value')

for ax in axs.flat:


```

```

ax.label_outer()
plt.show()

with open("Errors-S.csv", "r") as fp:
    x = from_csv(fp)
print(x)

```



Features		Test MAE	Test RMSE
Train MAE	Train RMSE		
Genre, Crew	0.7599656357	0.6829436039	0.8827476799
Crew, Keywords	0.9125429553	0.8418156809	1.036022444
Genre, Crew, Production Countries	0.776975945	0.7138927098	0.9162599656
Genre, Production Countries, Keywords	0.7374570447	0.6777166437	0.8672158173
Genre, Crew, Production Countries, Keywords	0.7451890034	0.6905089409	0.8905587631

```

+-----+-----+-----+-----+
-----+-----+

```

RMSE value close to 0 indicates that the regression model is able to generalise well. We take the set of attributes for which we get the lowest RMSE and MAE Errors. This our case, the required set of attriutes is: Genre, Production Countries, Keywords

We can also see that Test RMSE and Train RMSE are almost equal. Thus we can conclude that our model doesn't overfit/ underfit on the data.

Choosing Best Set of features according to minimum error

```
[0]: mmin = 1
      rmin = 1
      for i in range(len(testrmse)):
          if(mmin > testmae[i] and rmin > testrmse[i]):
              mmin = testmae[i]
              rmin = testrmse[i]
```

```
[0]: print(mmin,rmin)
```

```
0.6777166437414031 0.8672158172725218
```

```
[0]: bestFeat = feat[testmae.index(mmin)]
      print(bestFeat)
```

```
['genres', 'production_countries', 'keywords']
```

```
[0]: bestFeatVector = vectorForm(finalDf,bestFeat)
      bestFeatVector.to_csv('BestFeatureVector.csv',index=None) #Dumping best_
      ↪feature vector based on min MAE and RMSE to csv
```

Cosine Similarity

```
[0]: bfv = pd.read_csv('BestFeatureVector.csv') #Dumping best feature vector in csv
      simMatrix = cosine_similarity(bfv.values,bfv.values)
      pd.DataFrame(simMatrix).to_csv('cosSimScore.csv',index=None) #Dumping_
      ↪similarity matrix in csv file
      simMatrix = pd.read_csv('cosSimScore.csv')
      # simMatrix.shape
      sim = simMatrix.values
      mu = round(bfv.rating.mean(),2) #Mean rating for user who has done max ratings
      X = bfv.iloc[:,bfv.columns!='rating']
      y = bfv.rating
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
      ↪2,shuffle=False)
```

Finding Recommendations

```
[0]: recommend(X_test,y_test,sim,mu,finalDf.title,10)
```

```
[0]: ['Due cuori, una cappella',  
      'Samay: When Time Strikes',  
      'Virasat',  
      'The Last Relic',  
      'The Colour Out of Space',  
      'Shocking Asia',  
      'Koyla',  
      'The Incredible Kung Fu Master',  
      'The Mystery of Chess Boxing',  
      'Afterparty']
```